

Programming Knowledge with Frames and Logic

Michael Kifer

Stony Brook University

Part1: Foundations

What's in This Tutorial?

Part 1: Foundations

1. Introduction
2. Background
 - F-logic (Frame Logic)
 - HiLog
 - Transaction Logic
 - Top-down execution and tabling

What's in This Tutorial?

Part 2: Programming

3. Getting Around FLORA-2

- Getting started
- Modules
- Multifile modules
- Debugging

4. Some Low-level Details

- **HiLog** vs. Prolog representation of terms
- To table or not to table?

What's in This Tutorial?

5. Advanced Features

- Path expressions
- Aggregates
- Anonymous OIDs
- Equality
- Control constructs
- Metaprogramming

6. Updating the Knowledge Base

- Non-logical updates
- Logical updates
- Limitations
- Inserting and deleting rules

7. Future plans

1. Introduction

What's Wrong with Knowledge Representation Based on Classical Logic?

- Essentially flat data structures:
`person(John, '123 Main St.', 34)`
- Awkward meta-programming:
Which predicates mention John?
- Ill-suited for modeling side effects:
State changes, I/O

A Solution

- Flat data structures:

Frames (F-logic)

- Awkward meta-programming:

Higher-order syntax (HiLog + F-logic)

- Modeling side effects:

Logic of updates (Transaction Logic)

What is FLORA-2 ?

- **F-Logic tRAnslator**
- Realizes the vision of logic-based KR with frames, meta, and side-effects. Founded on
 - F-logic
 - HiLog
 - Transaction Logic
- Practical & usable KR and programming environment
 - Declarative
 - Object-oriented
 - Logic-programming style
 - Overcomes most of the usability problems with Prolog

What is FLORA-2?

- Builds on earlier experience with implementations of F-logic:
 - FLORID, FLIP, FLORA-1 (which don't support HiLog & Transaction Logic)
- Differs in spirit from other F-logic based systems
 - FLORID, Ontobroker are *query languages*; cannot live without a procedural language (C++, Java)
 - FLORA-2 is a complete *programming language*; can be used in the query language capacity as well.
- <http://flora.sourceforge.net>
- A recent overview: [Yang, Kifer, Zhao, ODBASE-2003]

Applications of FLORA-2

- Ontology management
- Knowledge-based networking
- Information integration
- Software engineering
- Agents
- Anything that requires manipulation of complex structured (especially semi-structured) data

Other F-logic Based Systems

- ????? (U. Melbourne – M. Lawley) – early 90's; first Prolog-based implementation
- *FLORID* (U. Freiburg – Lausen et al.) – late 90's; the only C++ based implementation
- *FLIP* (U. Freiburg – Ludaescher) – late 90's; first XSB based implementation. Inspired the FLORA effort
- *TFL* (Tech. U. Valencia – Carsi) – late 90's; first attempt at F-logic + Transaction Logic
- *SILRI* (Karlsruhe – Decker et al.) – late 90's; Java based
- *TRIPLE* (Stanford – Decker et al.) – early 2000's; Java
- *OntoBroker* (Ontoprise.de, now Semafora) – 2000; commercial

2. Background

Desirable Background Knowledge

- Predicate calculus
 - Good understanding of its model theory
- Logic programming/Deductive databases
 - Bottom-up execution (T_P operator)
 - Top-down execution (SLD resolution)
 - Negation as failure / Well-founded negation
- Prolog language

2.1. Background: F-Logic

Basic Ideas Behind F-logic

- Take complex data types as in object-oriented databases
- Combine them with logic
- Use the result as a programming language

What F-Logic Provides

- Objects with complex internal structure
- Class hierarchies and inheritance
- Typing
- Encapsulation
- Background:
 - Basic theory: [Kifer & Lausen SIGMOD-89], [Kifer,Lausen,Wu JACM-95]
 - Path expression syntax: [Frohn, Lausen, Uphoff VLDB-84]
 - Semantics for non-monotonic inheritance: [Yang & Kifer, ODBASE 2002]
 - Meta-programming + other extensions: [Yang & Kifer, ODBASE 2002]

Relationship to Standard Logic

O-O programming

Relational programming

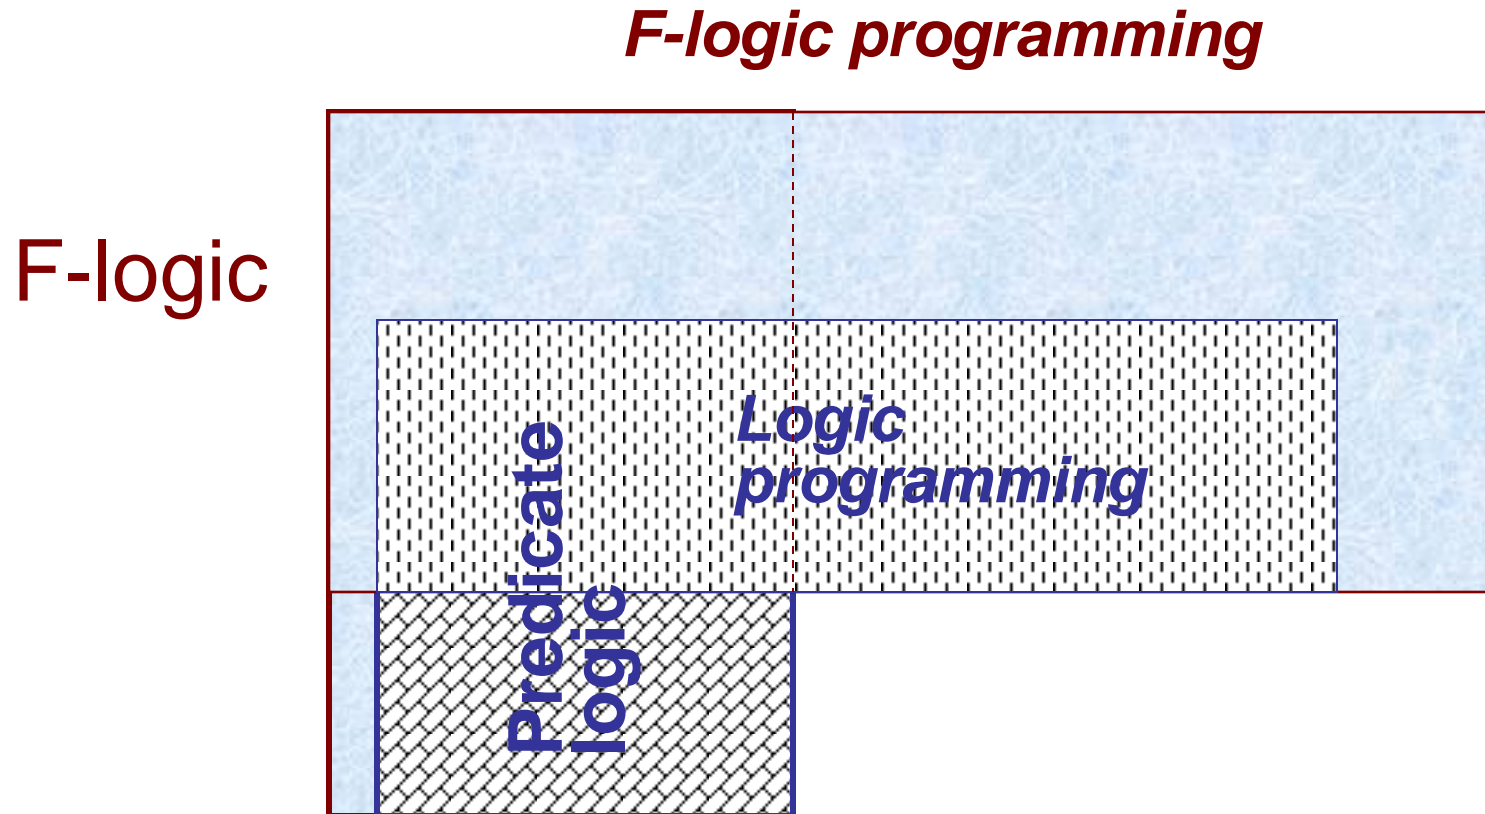
=

F-logic

Predicate calculus

Relationship to Standard Logic (cont'd)

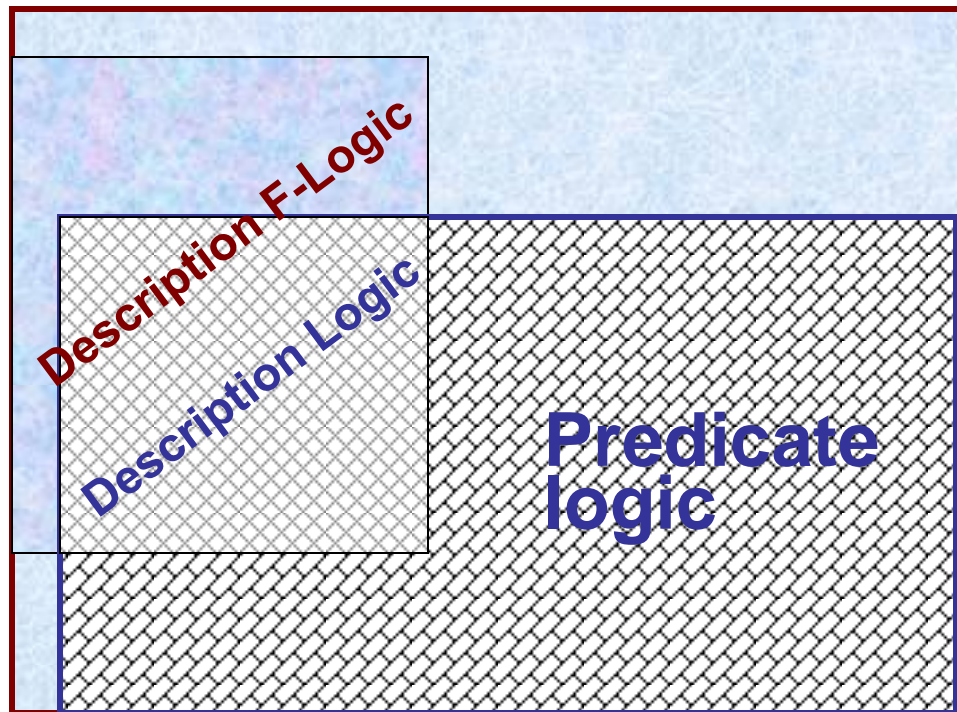
First-order flavor vs. logic programming flavor.



Relationship to Description Logic

A description logic subset can be developed in F-logic

[Balaban 1995, The F-logic Approach for Description Languages]



F-logic: Simple Examples

Object Id

Attribute

Object description:

John[*name* -> 'John Doe', *phones* -> {6313214567, 6313214566},
children -> {Bob, Mary}]

Mary[*name* -> 'Mary Doe', *phones* -> {2121234567, 2121237645},
children -> {Anne, Alice}]

Attribute

Structure can be nested:

Sally[*spouse* -> John[*address* -> '123 Main St.']]

Examples (cont'd)

- Historic notes:
 - The original F-logic distinguished between functional (\rightarrow) and set-valued ($\rightarrow\rightarrow$) attributes
 - In FLORA-2 this has been simplified and generalized:
 - Only set-valued methods and only \rightarrow are used
 - Can specify cardinality constraints. The constraint $\{0:1\}$ corresponds to functional attributes
 - In F-logic, variables were denoted by capitalized symbols
 - In FLORA-2 variables are preceded with a ?.
 - Constants can start with lowercase or uppercase – does not matter:
 - John, betty.

Examples (contd.)

ISA hierarchy:

John : Person - *class membership*

Mary : Person

alice : Student

Student :: Person - *subclass relationship*

Student : EntityType

Person : EntityType

Class & instance
at the same time

Examples (Contd.)

Methods: like attributes, but take arguments

$?P[\textit{ageAsOf}(?Year) \rightarrow ?Age] : -$

$?P:\textit{Person}, ?P[\textit{born} \rightarrow ?B], ?Age \text{ \textit{is} } ?Year - ?B.$

- Attributes can be viewed as methods with no arguments

Query :

John's children who were born when he was 30+ years old:

$?- \text{ John}[\textit{born} \rightarrow ?Y, \textit{children} \rightarrow ?C],$
 $?C[\textit{born} \rightarrow ?B], ?B > ?Y + 30.$

or

$?- \text{ John}[\textit{ageAsOf}(?Y) \rightarrow 30, \textit{children} \rightarrow ?C],$
 $?C[\textit{born} \rightarrow ?B], ?B > ?Y.$

Examples (Contd.)

- **Type signatures:** Define the types for method arguments and for their results

```
Person[born => integer,  
       ageAsOf(integer) => integer,  
       name => string,  
       address => string,  
       children => person].
```

- Signatures can be queried:
 ?– Person[*name* => ?Type].

Answer: ?Type = string

 ?– Person[?Attr => string].

Answer: ?Attr = name

 ?Attr = address

Syntax

- Object ids:
 - Terms like in Prolog, but constants, functions can be capitalized – John, abc, f(john,34), Car(red,20000)
 - Below, O, C, M, T, ... denote usual first order terms
- IsA hierarchy (*isa-atoms*):
 - O:C -- object O is a *member* of class C
 - C::S -- C is a *subclass* of S
- Structure (*object-atoms*):
 - O [Method → Value] -- invocation of method
- Type (*signature-atoms*):
 - Class [Method => Class] – a method signature
- Combinations of the above:
 - and, or, negation, quantifiers

More Examples

Browsing IsA hierarchy:

? – John : ?X.

? – Student :: ?Y

and

Virtual (view) class:

?X : Redcar : – ?X:Car , ?X[*color* -> red].

Meta-query about schema:

?O[*attributesOf*(?Class) -> ?A] : –
?O[?A -> ?V], ?V: ?Class.

*Rule defines method, which
returns attributes whose
range is class Class
 $\alpha :- \beta$ is implication, α or $\neg\beta$*

Parameterized family of classes:

[]:list(?T).

[?X|?L]:list(?T) : – ?X:?T, ?L:list(?T).

E.g., list(integer), list(student)

Model Theory for Object Definitions

Simplified (so-called *Herbrand*) semantics:

Universe: HB – set of all variable-free terms (“ground” terms)

Interpretation: $\mathbf{I} = (\text{HB}, I_{\rightarrow}, \in, <)$

where $< : \text{partial order on HB}$

$\in : \text{binary relationship on HB}$

$I_{\rightarrow} : \text{HB} \rightarrow (\text{HB} \xrightarrow{\text{partial}} \text{powerset}(\text{HB}))$

Satisfaction of formulas in \mathbf{I} :

$\mathbf{I} \models o[m \rightarrow v] \quad \text{if } v \in I_{\rightarrow}(m)(o)$

$\mathbf{I} \models o:c \quad \text{if } o \in c$

$\mathbf{I} \models c:s \quad \text{if } c < s$

methods

values

objects

Model Theory for Types

Interpretation: $I = (HB, I_{\rightarrow}, \in, <, I_{\Rightarrow})$

Added

where $I_{\Rightarrow} : HB \rightarrow (HB \xrightarrow{\text{partial}} \text{powerset}(HB))$

*The function assigns
types to methods*

set of methods

set of classes

types for results

Satisfaction of method signatures:

$I \models c[m \Rightarrow t]$ if *some* element in $I_{\Rightarrow}(m)(c)$ is $\leq t$

- Basically, we want $c[m \Rightarrow t]$ and $t :: t'$ to imply $c[m \Rightarrow t']$
(if the result is of type t then it also conforms to any supertype of t)

Semantics (cont'd)

The well-typing condition:

$o[m \rightarrow v]$ is *well-typed* in \mathbf{I}

iff whenever $o \in c$ then $v \in (I_{\rightarrow}(m)(c))$

\mathbf{I} is *well-typed* if every true object atom is well-typed.

Here we want $c[m \Rightarrow t]$, $o[m \rightarrow v]$, $o:c$ to imply $v:t$.

I.e., typing is a constraint

Semantics (cont'd)

- $\mathbf{I} \models P \wedge Q$ iff $\mathbf{I} \models P$ and $\mathbf{I} \models Q$
- $\mathbf{I} \models P \vee Q$ iff $\mathbf{I} \models P$ or $\mathbf{I} \models Q$
- $\mathbf{I} \models \neg P$ iff not $\mathbf{I} \models P$
- $\mathbf{I} \models \forall ?X P$ iff for all $c \in \text{HB}$, $\mathbf{I} \models P'$
 P' is P with *all* free occurrences of $?X$ replaced with c
- $\mathbf{I} \models \exists ?X P$ iff for some $c \in \text{HB}$, $\mathbf{I} \models P'$
 P' is P with *some* free occurrence of $?X$ replaced with c

Shorthands

- \wedge -Composition: $\mathbf{O}[m1 \rightarrow v1, m2 \rightarrow v2]$ is
$$\mathbf{O}[m1 \rightarrow v1] \wedge \mathbf{O}[m2 \rightarrow v2]$$
- \vee -Composition: $\mathbf{O}[m1 \rightarrow v; m2 \rightarrow v2]$ is
$$\mathbf{O}[m1 \rightarrow v1] \vee \mathbf{O}[m2 \rightarrow v2]$$
- *Nesting*: $\mathbf{O}[m1 \rightarrow \mathbf{v1}[m2 \rightarrow v2]]$ is
$$\mathbf{O}[m1 \rightarrow \mathbf{v1}] \wedge \mathbf{v1}[m2 \rightarrow v2]$$
- IsA-Composition: $\mathbf{O}:\mathbf{C}[m \rightarrow v]$ (or $\mathbf{O}[m \rightarrow v]:\mathbf{C}$) is
$$\mathbf{O}:\mathbf{C} \wedge \mathbf{O}[m \rightarrow v]$$
- Same for the other arrows

These are
called
molecules
or *frames*

Boolean Methods

- Another shorthand: `Obj[Meth]`
 - E.g. `?X[p(a,?X)]`, `f(?X)[p]`, `john[married(1999)]`
- Think of these as a shorthand for `Obj[Meth -> void]`

(this is only conceptually: `Obj[Meth]` is an independent construct and is not equivalent to `Obj[Meth -> void]`)
- **Boolean signatures:** `Obj[=>MethType]`
 - E.g., `Person[=>married(Year)]`

Proof Theory

- Resolution-based
 - Will see later a special case
- Sound & complete w.r.t. the semantics
 - Soundness of proofs:

If can prove Q from a set of formulas \mathbf{P} then $\mathbf{P} \models Q$
 - Completeness of proofs:

If $\mathbf{P} \models Q$ then can prove Q from \mathbf{P}

A Note on the Semantics of FLORA-2

- F-logic semantics & proof theory is completely general, like that of classical logic
- But FLORA-2 is a programming language, hence it uses non-classical semantics

$\dots :- \dots, \backslash\text{naf } P, \dots$

means: *true if cannot prove P* – so called “negation as failure.”

The exact semantics for negation used in FLORA-2 is Van Gelder’s Well-Founded Semantics [Van Gelder et al., JACM 1991, <http://citeseer.nj.nec.com/gelder91wellfounded.html>]

A Note on the Semantics (cont'd)

- The Well-Founded semantics is *3-valued*:

$p :- \backslash\text{naf } q.$

$r :- \backslash\text{naf } r.$

p is true, q false, but r is undefined

- And *non-monotonic*:

$P \models Q$ doesn't imply $P \cup P' \models Q$

$p :- \backslash\text{naf } q$ implies p true.

But


q and $p :- \backslash\text{naf } q$ implies p false.

- Classical logic is both *2-valued* and *monotonic*

Inheritance in Flora-2

- Inheritance of *structure* vs. inheritance of *behavior*
 - ***Structural inheritance*** = inheritance of the signature of a method
 - ***Behavioral inheritance*** = inheritance of the definition of a method
- Attributes/methods can be *class-level* and *object-level*
 - ***Object-level*** statements about an object, **c**, which may be a class-object, apply only to **c** and nothing else
 - ***Class-level*** statements are *inherited* from **c**. That is, they apply to all members of the class **c** and to all subclasses of **c**.

Structural Inheritance

- Class-level signatures appear inside class-level statements (`[|...|]`). Object-level signatures appear inside object-level statements (`[...]`).
 - For *object-level* statements:
 - `class[method => type]` and `subclass::class`
does **not** imply `subclass[method => type]`
 - For *class-level* statements:
 - `class[|method => type|]` and `subclass::class`
does **imply** `subclass[|method => type|]`
 - `class[|method => type|]` and `obj::class`
does **imply** `obj[method => type]`
 - *Structural inheritance is monotonic*: adding more signatures doesn't invalidate old inferences
- 

Structural Inheritance - Semantics

Interpretation: $I = (HB, I_{\rightarrow}, \in, <, I_{\Rightarrow}, I_{[| \Rightarrow |]})$

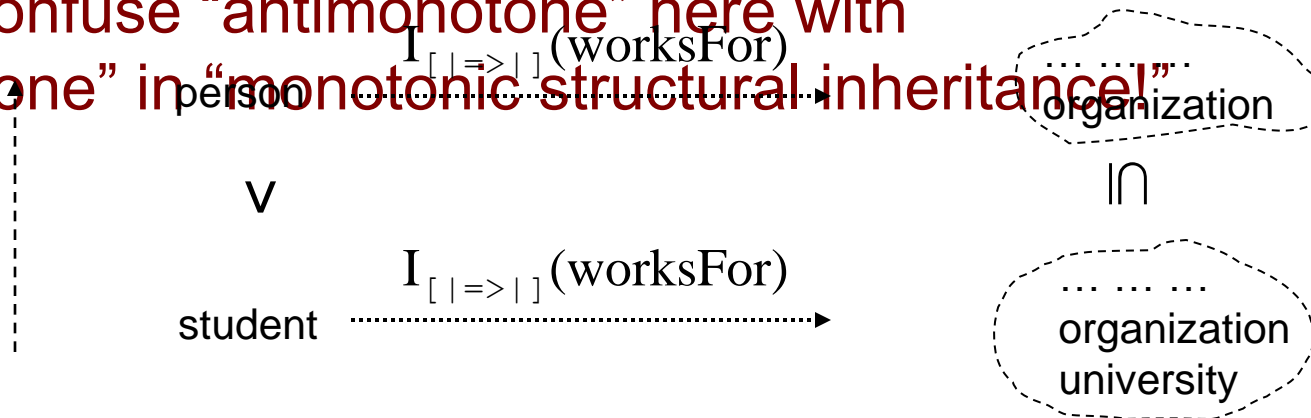
Added

where

$I_{[| \Rightarrow |]}: HB \rightarrow (HB \xrightarrow{\text{partial and antimonotone}} \text{powerset}(HB))$

Why antimonotonicity?



Don't confuse "antimonotone" here with "monotone" in "monotonic structural inheritance!"



Behavioral Inheritance


- Class-level statements use $\dots [\mid \dots - > \dots \mid]$
 - Object-level statements use $\dots [\dots - > \dots]$
- Behavioral inheritance is *non-monotonic*

Relationship Between Inheritable and Non-inheritable Methods




Inheritable methods are inherited as

- *inheritable* to subclasses
- *non-inheritable* to members



Behavioral Inheritance: Non-monotonicity



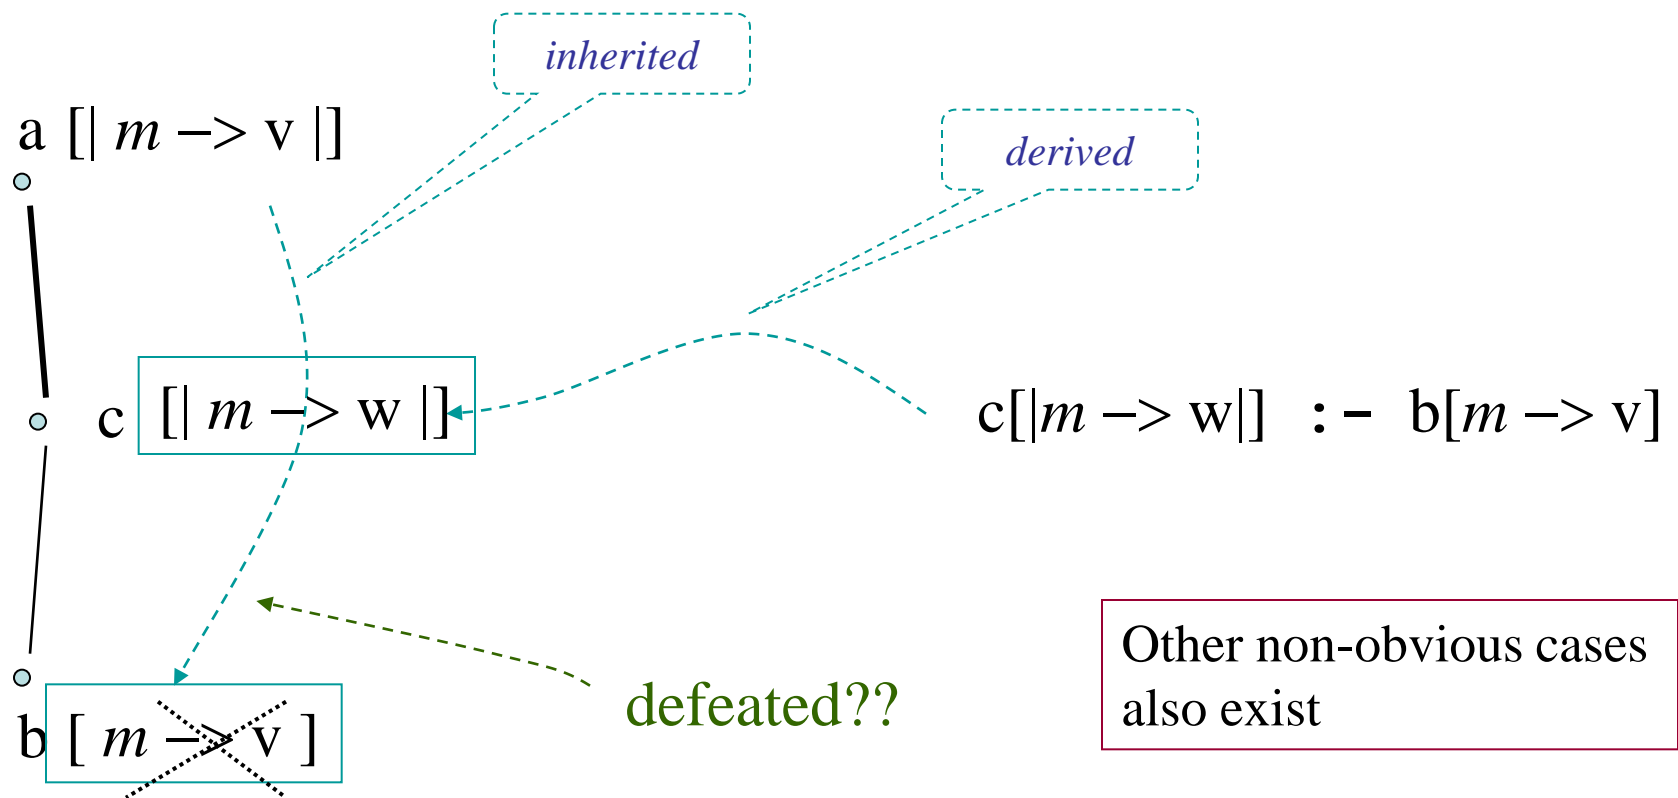
Inherited:

fred[*color* -> grey]
clyde[*color* -> ~~grey~~ ^{white} ?]

Overriding

Behavioral Inheritance: Problem with Rules

- Inheritance is hard to even define properly in the presence of rules.



Behavioural Inheritance: Solutions

- Hard to define semantics for multiple inheritance + overriding + rules
 - Several semantics might look “reasonable”
 - Should have no unnecessary restrictions
- The original semantics in [Kifer,Lausen,Wu: JACM-95] was one of the problematic “reasonable” semantics
 - A number of other problematic semantics of various degrees of “reasonableness” exist
- Problem solved in [Yang&Kifer: Journal on Data Semantics 2006]
 - Based on semantic postulates
 - An extension of Van Gelder’s Well-Founded Semantics for negation

2.2. Background: HiLog

HiLog

- Allows certain forms of logically clean meta-programming
- Syntactically appears to be higher-order, but semantically is first-order and tractable
- Has sound and complete proof theory
- [Chen,Kifer,Warren, HiLog: A Foundation for Higher-Order Logic Programming, J. of Logic Programming, 1993]
 - The recent work on SKIF and Common Logic (Hayes et. al.) is a rediscovery of **HiLog** with very minor differences – 12 years later!

Examples of HiLog

Variables over predicates and function symbols:

$p(?X, ?Y) : - ?X(a, ?Z), ?Y(?Z(b)).$

Variables over atomic formulas (*reification*):

$call(?X) : - ?X.$

A use of HiLog in FLORA-2 (e.g., querying of schema):

$?O[unaryMethods(?Class) \rightarrow ?M] : -$
 $?O[?M(?) \rightarrow ?V], ?V : ?Class.$

*Meta-variable: ranges
over method names*

Syntax and Semantics of HiLog

- In predicate logic, predicates and functions are disjoint, but predicate expressions (*atomic formulas*) and functional expressions (*function terms*) have the same syntax: e.g., $p(?X, f(a,b))$ vs. $g(?X, f(a,b))$
- **HiLog** makes no distinction between predicates and function symbols and atomic formulas are indistinguishable from function terms

Syntax of HiLog

- Everything is built out of constant symbols and variables
- **HiLog term:**
 - $?X$ and f (if $?X$ is a variable, f – a constant)
 - $F(A_1, \dots, A_n)$ if F, A_1, \dots, A_n are HiLog terms
 - Note: these are HiLog terms
 - Any Prolog term is, of course, a HiLog term
 - $X(a, f(?Y)), f(f(f, g), ?Y(?Y, ?Y)), h, ?Y$
 - $?X(a, f(Y))(f(f(f, g), Y(Y, Y)), h, Y)$
 - $?X(a, f(?Y))(X(a, f(?Y)))(f(f(f, g), ?Y(?Y, ?Y)), h, ?Y)$
- **HiLog formula:**
 - Any HiLog term
 - $A \vee B, A \wedge B, \neg A, \forall X A$, etc., if A, B are Hilog formulas

The “weird”
ones

Syntax of HiLog:

What are the “Weird” terms for?

- Generic transitive closure:

$\text{transClosure}(\text{?P})(\text{?X}, \text{?Y}) :- \text{?P}(\text{?X}, \text{?Y}).$

$\text{transClosure}(\text{?P})(\text{?X}, \text{?Y}) :- \text{?P}(\text{?X}, \text{?Z}), \text{transClosure}(\text{?P})(\text{?Z}, \text{?Y}).$

- For instance:
 - $\text{transClosure}(\text{parent})$ is the ancestor relation
 - $\text{transClosure}(\text{edge})$ pairs of all reachable nodes in the graph defined by *edge*

Semantics of HiLog

- **Interpretation** (Herbrand, for simplicity):
 - \mathbf{I} = any set of variable-free HiLog terms
 - $\mathbf{I} \models a$ (atomic variable-free), if $a \in \mathbf{I}$
 - $\mathbf{I} \models \phi \wedge \psi$, if $\mathbf{I} \models \phi$ and $\mathbf{I} \models \psi$
 - etc. (as usual)
 - $\mathbf{I} \models \forall X \phi$, if for all constant symbols c , $\mathbf{I} \models \phi[X \backslash c]$, where $\phi[X \backslash c]$ is ϕ with free occurrences of X replaced with c

Relationship to Predicate Logic

- $\models_{\text{classical}} \psi$ implies $\models_{\text{hilog}} \psi$
- $\models_{\text{hilog}} \psi$ does **not** imply $\models_{\text{classical}} \psi$:
 - $(q(a) \leftrightarrow r(a)) \leftarrow \forall X \forall Y (X=Y)$
is valid in HiLog but not in predicate logic
- But:
 - $\models_{\text{hilog}} \psi$ implies $\models_{\text{classical}} \psi$, except for formulas that are true in every interpretation with at least γ elements in the domain (for some $\gamma > 0$), but are false in some interpretation that has less than γ elements [Chen, Kifer, Warren JLP-93].
 - Examples: Horn clauses without “=” in the head;
Any set of “=”-free formulas

Reification:

An Application of HiLog to F-logic

- **Reification:** makes an object out of a statement:

john[believes -> **$\{$ mary[likes -> bob $\}$**]

- Introduced in [Yang & Kifer, ODBASE 2002]

- **Main idea:**

- Extend the syntax of F-logic to allow terms of the form

$\{$ mary[likes -> bob $\}$, **$\{$ bob[name -> 'Bob Doe' $\}$**

and even more general ones, like

$\{$ mary[likes -> bob, name -> 'Bob Doe' $\}$

- Eliminate the distinction between atomic formulas and terms both in the syntax and semantics (like in HiLog)

Object made out of
the statement
mary[likes -> bob]

The Role of HiLog

- HiLog and its applications to F-logic (*reification*, *schema browsing*) allows high degree of meta-programming purely in logic
- Variables can be bound to predicate and function symbols and thus queried (e.g., which relation mentions constant ‘john’)
- Formulas can be represented as terms, decomposed, composed, and manipulated with in flexible ways
- One can mix frame syntax (F-logic) and predicate syntax (HiLog) in the same query/program:
a[b \rightarrow c, g(?X,e) \rightarrow d], p(f(?X),a).

2.3. Background: Transaction Logic

Transaction Logic

- A logic of change
- Unlike temporal/dynamic/process logics, it is also a logic for *programming* (but can be used for *reasoning* as well)
- In the object-oriented context:
 - A logic-based language for programming the *behavior* of objects, i.e., specifying methods that change the object state

[Bonner&Kifer, An Overview of Transaction Logic, in *Theoretical Computer Science*, 1995],

[Bonner&Kifer, A Logic for Programming Database Transactions, in *Logics for Databases and Information Systems*, Chomicki+Saake (eds), Kluwer, 1998].

[Bonner&Kifer, Results on Reasoning about Action in Transaction Logic, in *Transactions and Change in Logic Databases*, LNCS 1472, 1998].

What's Wrong with Other Logics for Specifying Change?

- Designed for reasoning, *not* programming
 - E.g., situation calculus, temporal, dynamic, process logics
- Typically lack such basic facility as subroutines
- None became the basis for a reasonably useful programming language

Problems with Specifying Change in Logic Programming (Prolog)?

- *assert/retract* have no logical semantics
- Non-backtrackable, e.g.,

?– $\text{assert}(p), q.$

If q is false, p stays.

- Prolog programs with updates are the hardest to write, debug, and understand

Example: Stacking a Pyramid

Program:

stack(0,X).

stack(N,X) :- N>0, move(Y,X), stack(N-1,Y).

move(X,Y) :- pickup(X), putdown(X,Y).

pickup(X) :- clear(X), on(X,Y), retract(on(X,Y)), assert(clear(Y)).

putdown(X,Y) :- wider(Y,X), clear(Y), assert(on(X,Y)), retract(clear(Y)).

Action:

?- *stack(18,block32).* // stack 18-block pyramid on top of block 32

Note:

Prolog *won't* execute this intuitively correct program properly!

Syntax

- Serial conjunction, \otimes (often denoted using “,”)
 - $a \otimes b$ – do a then do b
- The usual $\wedge, \vee, \neg, \forall, \exists$ (but with a different semantics)
 - Example: $a \vee (b \otimes c) \wedge (d \vee \neg e)$
- $a :- b \equiv a \vee \neg b$
 - Means: to execute a one must execute b (i.e., a is a subroutine)
- Transaction logic also has hypothetical operators \diamond and \Box , but won't discuss (not implemented in FLORA-2)

Semantics

- Model-theoretic, like F-logic and HiLog
 - The basic ideas
 - *Execution path* \equiv sequence of database states
 - Assume that the states are just sets of facts
 - Truth values over paths, not over states
 - Truth over a path \equiv *execution* over that path
 - *Elementary state transitions* \equiv propositions that cause a priori defined state transitions
 - For most purposes, can use the following elementary state transitions:
 - $t_insert\{fact\}$ and $t_delete\{fact\}$ (for *transactional* insert and delete)
 - $t_insert\{fact\}: D \rightarrow D + fact$ – add *fact* to state **D**
 - $t_delete\{fact\}: D \rightarrow D - fact$ – delete *fact* from state **D**
 - FLORA-2 allows more powerful state transitions (**bulk updates**):
 - $t_insert\{fact(?X)|condition(?X)\}$ and $t_delete\{fact(?X)|condition(?X)\}$
- Insert/delete things of the form $fact(X)$ that satisfy $condition(X)$.

Path Structures

- Semantics is defined using the notion of path structures (which play the same role as semantic structures in classical logic)
- A *path structure* maps execution paths to the ordinary semantic structures used in classical predicate logic:

$I(\pi) = M$, where π - path, M – classical semantic structure, which says which transactions can execute along the path π


In addition:

- If $\pi = \langle D \rangle$ is a path that consists of only one database state then $I(\pi)$ must make every fact in D true.
- If $\pi = \langle D, D + \text{fact} \rangle$ then $I(\pi)$ should make $t_insert\{\text{fact}\}$ true
- If $\pi = \langle D, D - \text{fact} \rangle$ then $I(\pi)$ should make $t_delete\{\text{fact}\}$ true

Satisfaction

Intuition:

$a \otimes b$: First execute a then b – represents sequencing of actions




Definition:

$\mathbf{I}(\langle D_0, \dots, D_n \rangle) \models a \otimes b$ iff $\exists D_k$ such that $\mathbf{I}(\langle D_0, \dots, D_k \rangle) \models a$ and $\mathbf{I}(\langle D_k, \dots, D_n \rangle) \models b$

Satisfaction (cont'd)

Intuition:

$a \wedge b$: Execute a along a path that is also an execution of b – represents constraints



Definition:

$I(<D_0, \dots, D_n>) \models a \wedge b$ iff $I(<D_0, \dots, D_n>) \models a$ and $I(<D_0, \dots, D_n>) \models b$


Satisfaction (cont'd)

Intuition:

$a \vee b$: Execute a along a path or execute b – represents choice

If: $I(\langle D_0, \dots, D_7 \rangle) \models a$

or: $I(\langle D_0, \dots, D_7 \rangle) \models b$



Then: $I(\langle D_0, \dots, D_7 \rangle) \models a \vee b$

Definition:


$I(\langle D_0, \dots, D_n \rangle) \models a \vee b$ iff $I(\langle D_0, \dots, D_n \rangle) \models a$ or $I(\langle D_0, \dots, D_n \rangle) \models b$

Satisfaction (cont'd)

Intuition:

$\neg a$: Execute in any way provided that it is not an execution of a

If: $I(<D_0, \dots, D_7>) \not\models a$



Then: $I(<D_0, \dots, D_7>) \models \neg a$

Definition:


$I(<D_0, \dots, D_n>) \models \neg a$ iff $I(<D_0, \dots, D_n>) \not\models a$

Satisfaction (cont'd)

head \leftarrow *body* (defined as $a \vee \neg b$)

Formally: Every execution of *body* is also an execution of the *head*:

If: $\mathbf{I}(\langle D_0, \dots, D_7 \rangle) \models \textit{body}$




Then: $\mathbf{I}(\langle D_0, \dots, D_7 \rangle) \models \textit{head}$

Informally: One way to execute *head* is to execute *body*
 \Rightarrow *head* is the name of a procedure
and *body* is part of its definition

Properties of the Semantics

The semantics has the “*all or nothing*” flavor which makes updates logical:



If action is *true*, but postcondition *false*, then
 $\text{action} \otimes \text{postcondition}$ is *false* on π .

In practical terms: *updates are undone on backtracking.*

Transaction Programs

- A *transaction program* P is a set of rules of the form *head* : − *body* like

move(?X,?Y) : − pickup(?X), putdown(?X,?Y)

which define complex transactions using simple actions (like *t_insert/t_delete*)

- A *transaction* (or action) is a query of the form
? − *body*.
(e.g., ? − *stack(18,block32)*)

Proof Theory

- *Executorial entailment*: \mathbf{P} is a set of rules, ϕ is a transaction (query), D_1, \dots, D_n — a sequence of states. Then

$$\mathbf{P}, D_1, \dots, D_n \models \phi$$

iff \forall path structures \mathbf{I} where $\mathbf{I} \models \mathbf{P}$ (ie., \forall path π , $\mathbf{I}(\pi) \models \mathbf{P}$),

it follows that $\mathbf{I}(\langle D_1, \dots, D_n \rangle) \models \phi$

- To prove ϕ from a set of rules (transaction definitions) \mathbf{P} , the proof theory tries to find a path, D_1, \dots, D_n , on which ϕ is executionally entailed by \mathbf{P} .
 - Thus, the proof theory *executes* ϕ as it proves it (and changes the underlying database state from the initial state D_1 to the final state D_n)

Pyramid Building (again)

stack(0, ?X).

stack(?N, ?X) :- ?N > 0 \otimes move(?Y, ?X) \otimes stack(?N-1, ?Y).

move(?X, ?Y) :- pickup(?X) \otimes putdown(?X, ?Y).

pickup(?X) :- clear(?X) \otimes on(?X, ?Y) \otimes t_delete{on(?X, ?Y)} \otimes t_insert{clear(?Y)}.

putdown(?X, ?Y) :- wider(?Y, ?X) \otimes clear(?Y) \otimes t_insert{on(?X, ?Y)} \otimes t_delete{clear(?Y)}.

?- stack(18, block32). // stack 18-block pyramid on top of block 32

- Under the Transaction Logic semantics the above program does the right thing

Constraints

- Can express not only execution, but all kinds of sophisticated constraints:

$?- \text{stack}(10, \text{block43})$

$\wedge \quad \forall ?X, ?Y (\text{move}(?X, ?Y) \otimes \text{color}(?X, \text{red})) \Rightarrow (\exists ?Z \text{color}(?Z, \text{blue}) \otimes \text{move}(?Z, ?X))$

Whenever a red block is stacked, the next block to be stacked must be blue

- Extensions (concurrent, game-theoretic) have been shown useful for process modeling
 - [Davulcu, Kifer, Ramakrishnan, & Ramakrishnan, Logic Based Modeling and Analysis of Workflows, in Proceedings of *PODS*, 1997]
 - [Davulcu, Kifer, Ramakrishnan, CTR-S: A Logic for Specifying Contracts in Semantic Web Services, Proceedings of WWW2004]

Reasoning

- Can be used to *reason* about the effects of actions such as:
 - If ϕ was true before the execution of transaction then ψ must be true after
 - If ϕ was true after the execution of transaction then ψ must have been true before

[Bonner&Kifer, Results on Reasoning about Action in Transaction Logic, in *Transactions and Change in Logic Databases*, LNCS 1472, 1998]

Planning

- Transaction Logic is ideal for specifying planning strategies.
- **The planning problem:**
 - *Given:*
 - A set of *primitive actions* – a_1, \dots, a_n
each a_i can have preconditions
 - A *goal* – G
a condition on the final state of the DB,
which we want to achieve
 - An *initial state* D_0
 - *Find:*
 - A sequence of these actions that starting at D_0 leads to a state D that satisfies G .

Naïve Planning is Easy in Transaction Logic

Specification:

$plan : - action \otimes plan.$

$plan : - action.$

$action : - a_1.$

... ..

$action : - a_n.$

To find a plan, just pose the query

$?- plan \otimes goal.$

Example:

$?- plan \otimes (on(b,c) \wedge on(c,d) \wedge clear(b)).$

Problem:

Proof theory might search through all sequences.

Planning with Heuristics

- Planning strategies employ heuristics to avoid exhaustive search
- **Transaction Logic** is ideal for specifying (and executing!) such heuristics
- Will illustrate using STRIPS (a classic planning system) as an example

STRIPS

- Uses actions of the form:

Name: *unstack(?X,?Y)*

Comment: Pick up block X from block Y

Precondition: *handempty, clear(?X), on(?X,?Y)*

Delete: *handempty, clear(?X), on(?X,?Y)*

Insert: *clear(?Y), holding(?X)*

- Uses an ad hoc algorithm to construct plans
- Most AI planning systems use ad hoc algorithms
- We can write planning strategies at the high level in **Transaction Logic** without worrying about the low-level details

Specifying STRIPS in Transaction Logic

- First, write a rule for each action – straightforward

$$\begin{aligned} \textit{unstack}(\textit{?X}, \textit{?Y}) : - & \textit{handempty} \otimes \textit{clear}(\textit{?X}) \otimes \textit{on}(\textit{?X}, \textit{?Y}) \\ & \otimes \textit{t_delete}\{\textit{clear}(\textit{?X}), \textit{on}(\textit{?X}, \textit{?Y}), \textit{handempty}\} \\ & \otimes \textit{t_insert}\{\textit{holding}(\textit{?X}), \textit{clear}(\textit{?Y})\} \end{aligned}$$

STRIPS in Transaction Logic (cont'd)

- Next, show how to *achieve* each goal of interest

achieve_clear(?Y) : - *achieve_unstack*(?X,?Y).

achieve_holding(?X) : - *achieve_unstack*(?X,?Y).

achieve_unstack(?X,?Y) : -

$(\textit{achieve_clear}(\textit{?X}) * \textit{achieve_on}(\textit{?X}, \textit{?Y}) * \textit{achieve_handempty})$

$\otimes \textit{unstack}(\textit{?X}, \textit{?Y}).$

(We use $a*b$ as a shorthand for $(a \otimes b) \vee (b \otimes a).$)

- The above says:
 - To achieve a goal, achieve the precondition of an action that inserts that goal
 - To achieve a precondition, achieve each of the subgoals in that precondition

STRIPS in Transaction Logic (cont'd)

- Base case: if a goal is already true, then it has been achieved

achieve_on(?X,?Y) : - on(?X,?Y).

achieve_clear(?X) : - clear(?X).

achieve_holding(?X) : - holding(?X).

achieve_handempty : - handempty.

STRIPS in Transaction Logic (cont'd)

- A STRIPS planning query in Transaction Logic
 - Stack c on d and b on c
 $?- (\text{achieve_on}(b,c) * \text{achieve_on}(c,d)) \otimes \text{on}(b,c) \otimes \text{on}(c,d).$
- The above is “ultimate” STRIPS: it finds a solution when one exists
- STRIPS was not based on a logic, so they kept refining their ad hoc execution mechanism
 - The original STRIPS was not complete. Was made complete after a series of papers
- The right logic makes the whole problem almost trivial!

Concurrent Transaction Logic

- Extends **Transaction Logic** with two connectives:
 - $a \mid b$ – *parallel conjunction*, denotes parallel execution
 - $\odot a$ – *isolation*, denotes isolated execution (in the sense of transaction processing)
 - Extends the model theory and the proof theory of **Transaction Logic**
- [Bonner&Kifer, Concurrency and Communication in Transaction Logic, in *Joint Int'l Conference and Symposium on Logic Programming*, MIT Press, 1996]
- Suitable for process modeling and programming concurrent systems
 - [Davulcu, Kifer, Ramakrishnan, & Ramakrishnan, Logic Based Modeling and Analysis of Workflows, in *Proceedings of PODS*, 1997]
- Harder to implement (not implemented in FLORA-2)
 - An interpreter available at <http://www.cs.toronto.edu/~bonner/ctr/>

Concurrent Transaction Logic for Services

- Extends **Concurrent Transaction Logic** with one additional connective:
 $a \sqcap b$ – the *opponent's conjunction*
- Enables specification of the *behavioral aspects* of *service contracts*
 - When different parties to the contract can make different choices (e.g., ship insured or uninsured, pay in full or in installments)
- [Davulcu, Kifer, & Ramakrishnan, **CTR-S**: A Logic for Specifying Contracts in Semantic Web Services, *WWW 2004*, May 2004]

2.4. Background: Top-down Execution and Tabling

SLD-Resolution

- Strategy at the core of any top-down execution engine
- *Sound* inference strategy
- *Complete* only for pure **Horn** clauses, i.e.,
 - Set of **rules**: $head : - body$ where *head* is atomic (of the form $p(\dots)$) and *body* is b_1, \dots, b_n (conjunction of atomic formulas). No negation in the head or the rule body.
 - Can be viewed as $head \vee \neg b_1 \vee \dots \vee \neg b_n$
 - Set of **facts**: atomic formulas.
 - Same syntax as *head*.
 - Can be viewed as a rule with empty body.
 - **Goal**: same syntax as the rule body.
 - The purpose of SLD resolution is to prove that $\exists ?X \text{ goal}$ ($?X$ represents all the vars in *goal*) follows from the set of facts plus the set of rules
 - Find all x such that $goal[?X \backslash x]$ (goal in which all occurrences of $?X$ are replaced with x) is implied by *rules* + *facts*.

SLD (cont'd)

- Goal: g_1, \dots, g_k

Rule: $h : - b_1, \dots, b_n$

☞ Rename vars in the rule to be disjoint from the vars in goal

θ : most general substitution s.t. $h\theta = g_1\theta$

- Derive new goal: $(b_1, \dots, b_n, g_2, \dots, g_k)\theta$

Note: g_1 replaced with b_1, \dots, b_n

- *Example:*

- Goal: $p(?X, f(?Y)), q(?X, ?Y, ?Z)$
- Rule: $p(g(?V), ?W) : - r(?V, f(?W)), h(?W, ?U).$
- θ : $?X \rightarrow g(?V), ?W \rightarrow f(?Y)$
- Derived goal: $r(?V, f(f(?Y))), h(f(?Y), ?U), q(g(?V), ?Y, ?Z)$

SLG (SLD with negation)

- When rules have negation in the body, the logically sound approach is to use the 3-valued Well-Founded Semantics (mentioned earlier)
- The adaptation of SLD to this case is called **SLG Resolution**. [Swift and Warren, *Intl. Logic Programming Symposium*, 1994]
 - Roughly works as SLD, but when it sees $\text{\texttt{\textbackslash naf } } p$ in the rule body, tries to prove p , possibly delaying until the literals to the right of $\text{\texttt{\textbackslash naf } } p$ have been proved. Three outcomes:
 - Proved p : $\text{\texttt{\textbackslash naf } } p$ is **false**
 - Proved that p cannot be proved: $\text{\texttt{\textbackslash naf } } p$ is **true**
 - All ways of deriving p rely on assuming $\text{\texttt{\textbackslash naf } } p$: p is **undefined**

Prolog Execution Strategy

- What if several rules have heads that unify with g_1 in g_1, \dots, g_k ?
 - SLD doesn't assume any order in which these rules are tried. If all orders are tried, then SLD is complete for Horn rules
 - Prolog does assume an order: rules are tried in the order in which they occur in the program. This causes Prolog to miss solutions even if they exist:
 - Goal: $?- p(?X)$
 - Rules: $p(?X) :- p(?X).$
 - $p(?X) :- r(?X).$
 - $r(a).$
- Prolog will get stuck in an infinite loop due to the first rule

Solution: Tabling

- When an attempt to solve a literal in the rule body is made (a *call* to the literal is made), save it in a table
- If the same call is made again, don't use SLD – look up the table instead; feed the answers from the first call to the second. Meanwhile, explore the other possibilities

- Example:

Goal: ?– p(?X)

Rules: p(?X) :– p(?X).

p(?X) :– r(?X).

r(a).

Call to p(?X). Save it in the table.

First derivation branch:

Use SLD with rule #1;

- create another call to p(?X).
- Look up the table—don't execute!
- Postpone this derivation branch.

Second derivation branch: Use SLD with rule #2

Call to r(?X). Save in the table.

Resolve with the fact r(a), get a result: ?X=a

No answers in the 1st derivation branch

Tabling (cont'd)

- See [Warren, CACM 1992]
- SLG resolution incorporates tabling
- SLG (unlike Prolog) is complete for Horn clauses; it is complete for the Well-Founded semantics for queries with negation in the rule body
- XSB is the only complete implementation of SLG
- YAP (<http://yap.sourceforge.net>) has an implementation of tabling; aims at having a complete implementation in the future

SLD and SLG in F-logic

- Similar to Prolog. Difference: goals and rule heads can have **F-logic** molecules in them:

Goal: ?- a[b -> c, d -> e].

Rules: ?Z[b -> ?Y, f -> ?Z] :- *body*.

 ?X[d -> ?Y, h -> ?Z] :- *anotherBody*.

Can these rules resolve with the goal?

- Answer: The notion of SLD resolution needs a slight modification.

SLD in F-logic (cont'd)

- Goals are transformed to eliminate disjunction (remember: disjunction is allowed in rule bodies and goals, but not in rule heads):

$?- ?X[\text{disj1} ; \text{disj2}], \text{rest}.$

becomes a pair of goals:

$?- ?X[\text{disj1}], \text{rest}.$

$?- ?X[\text{disj2}], \text{rest}.$

Must solve each goal and *union* the solutions.

- Note: a similar transformation is done in regular logic programming:

$?- (p ; q), \text{rest}.$

becomes

$?- p, \text{rest}.$

$?- q, \text{rest}.$

SLD in F-logic (cont'd)

- Goals are further transformed to simplify molecules:

$?- ?X[\text{part1}, \text{part2}], \text{rest}.$

becomes

$?- ?X[\text{part1}], ?X[\text{part2}], \text{rest}.$

and

$?- ?X[\text{foo} \rightarrow \{\text{bar1}, \text{bar2}\}], \text{rest}.$

becomes

$?- ?X[\text{foo} \rightarrow \text{bar1}], ?X[\text{foo} \rightarrow \text{bar2}], \text{rest}.$

Break molecules down into *atomic* (indivisible) ones.

SLD in F-logic (cont'd)

- SLD rule:

Goal: $?- \text{ subgoal-atomic-molecule}, \text{ rest}.$

Rule: $\text{ head-molecule} : - \text{ body}.$

☞ Rename vars in the rule to be disjoint from the vars in the goal

θ : most general unifier of $\text{ subgoal-atomic-molecule}$ *into* head-molecule , i.e, $\theta(\text{subgoal-atomic-molecule}) \subseteq \theta(\text{head-molecule})$

(\subseteq means both have the same object-term and the single component of $\text{ subgoal-atomic-molecule}$ inside the [...] is one of the components of head-molecule)

New goal: $?- \theta(\text{body}), \theta(\text{rest}).$

SLD in F-logic (cont'd)

- Example:
 - $?- f(?X,a)[m1 \rightarrow ?X, m2(?Y) \rightarrow b], p(?Y).$
 - $?V[?W \rightarrow c, m2(?V) \rightarrow b, m1 \rightarrow ?W] :- a[?V \rightarrow ?W].$
- Transform:
 - $?- f(?X,a)[m1 \rightarrow ?X], f(?X,a)[m2(?Y) \rightarrow b], p(?Y).$
- One unifier and new goal:
 - $\theta: ?V \rightarrow f(?X,a), ?W \rightarrow m1, ?X \rightarrow c$
 - $?- a[f(?X,a) \rightarrow m1], f(?X,a)[m2(?Y) \rightarrow b], p(f(?X,a)).$
- Another possibility:
 - $\theta: ?V \rightarrow f(?X,a), ?W \rightarrow ?X$
 - $?- a[f(?X,a) \rightarrow ?X], f(?X,a)[m2(?Y) \rightarrow b], p(f(?X,a)).$

SLG in F-logic

- FLORA-2 uses Prolog-like execution strategy
 - To be complete, it uses tabling
 - For negation in the rule body, it uses the Well-Founded Semantics and thus the SLG resolution
- To support inheritance, it uses an *extended* Well-Founded semantics, as mentioned earlier.
 - This is implemented by a translation into a Prolog program, which utilizes SLG resolution